

graphsummit

DUBLIN

neo4j

Building a Graph Solution

Using a real world digital twin data set

Marco De Luca & Gal Bello

Neo4j Field Engineering Team

Goal of the Today

**Demonstrate, with YOUR HELP :-), how to build a
Graph Solution using a Digital Twin Dataset**

Agenda

- 1. Logistics**
- 2. Introduction**
- 3. Use Case Explanation**
- 4. Building the solution**

BREAK (15min)

- 5. Discussion how to improve the Use Case further**
- 6. Final Q&A**
- 7. Done**

Logistics

First things first ...



Logistics

WIFI Access:	WIFI Name: NEO4J / Password: NEO4J2023
Restrooms:	Out of the room and to left
Chargers:	Power is under the table
Material of the workshop: (Type or scan)	https://github.com/neo4j-field/gsummit2023 

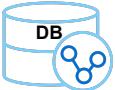
Introduction

A short overview of the Neo4j Product



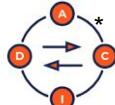
Neo4j Native Graph Database - What is it?

Neo4j is a NORMAL Database (DB)



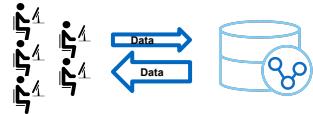
- Comparable with other DBs like Postgres, MySQL / MariaDB, Oracle, HANA DB, etc.
- Securely save, insert, update and query data
- Backup + Recovery
- Highly scalable

Neo4j is “ACID compliant” and transaction save



- **ACID (atomicity, consistency, isolation, durability)**
- Data is stored consistently on transaction level without any losses during an outage
- **Most relational DB's are ACID compliant, too.**

Neo4j is used 80%+ for OLTP Workloads



- > 80% of Neo4j Customers using the database as normal OLTP data store
- Data is stored, changed, deleted and queried
- With the same or even better performance compared to relational database systems

What is the Value Using a Native Graph-Database?

How data is stored on disk



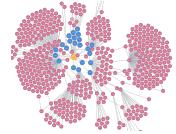
- Data is already stored connected
- Data is efficiently placed on disk using “Index-free Adjacency”
- A GDB stores **Nodes** and **Relations** instead of Rows and Columns
- **Semantics** may be build into the data!!

How to query the data

`(:Product) - [:CONTAINS] -> (:Part)`

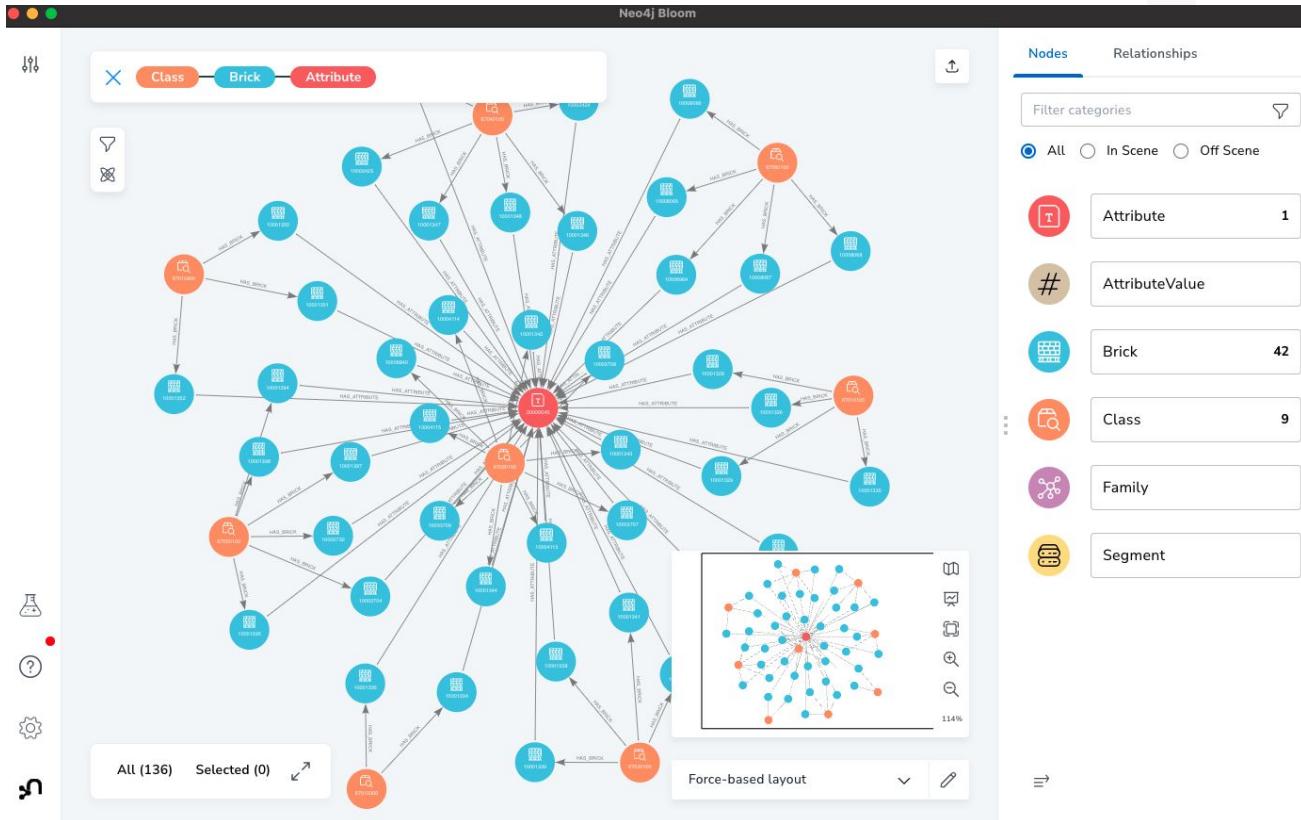
- Cypher Query Language vs. SQL (ISO -> GQL)
- Simple, lesser lines of code, easier to read and maintain
- Queries possible, that span 100+ Hops in the Graph, which is comparable with SQL Joins over 100+ Tables!

Storing complex data networks and semantics

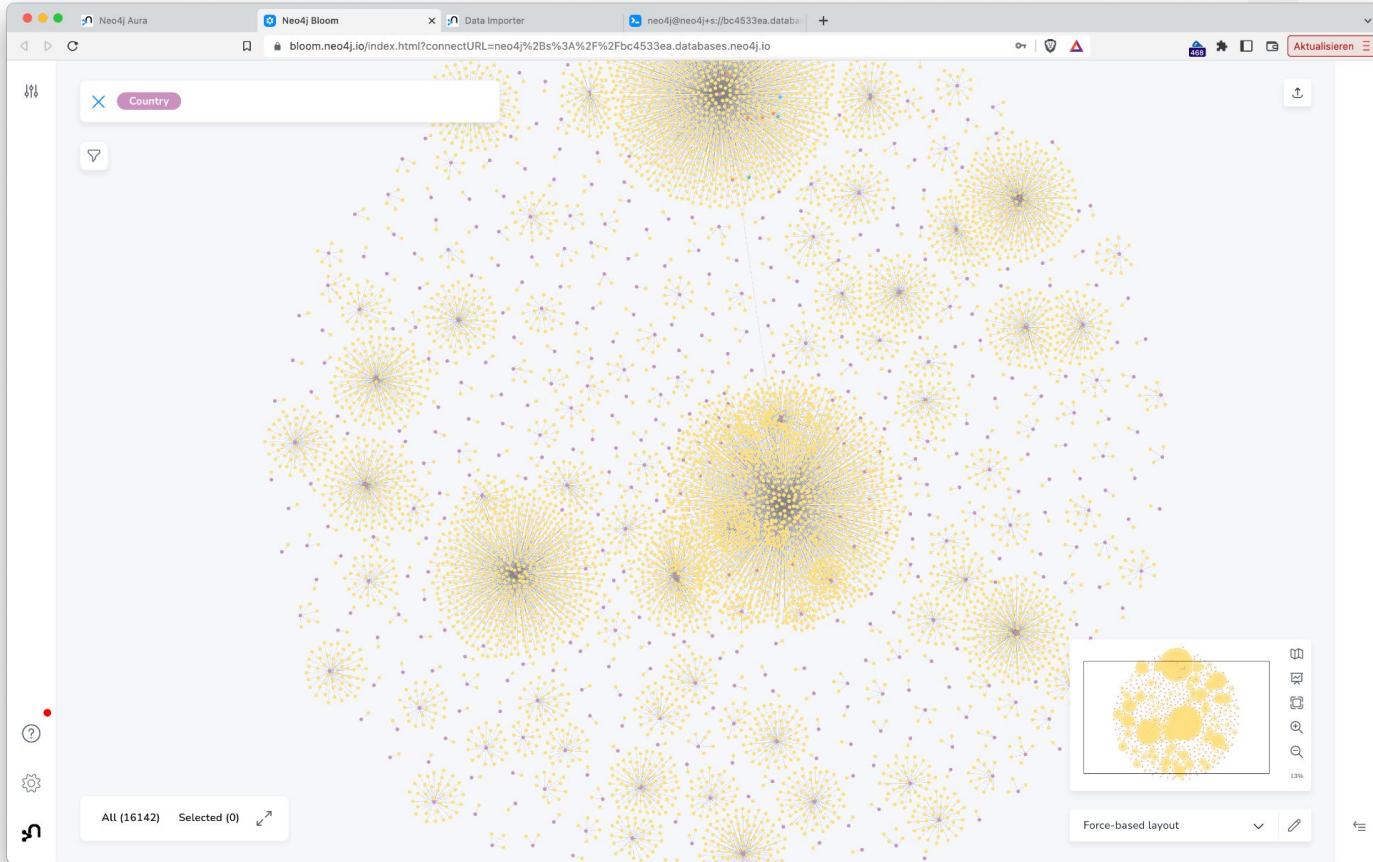


- Storing and analyzing complex relations between data
- Analyzing data from traditionally siloed data sources
- Extendable with data science algorithms and AI

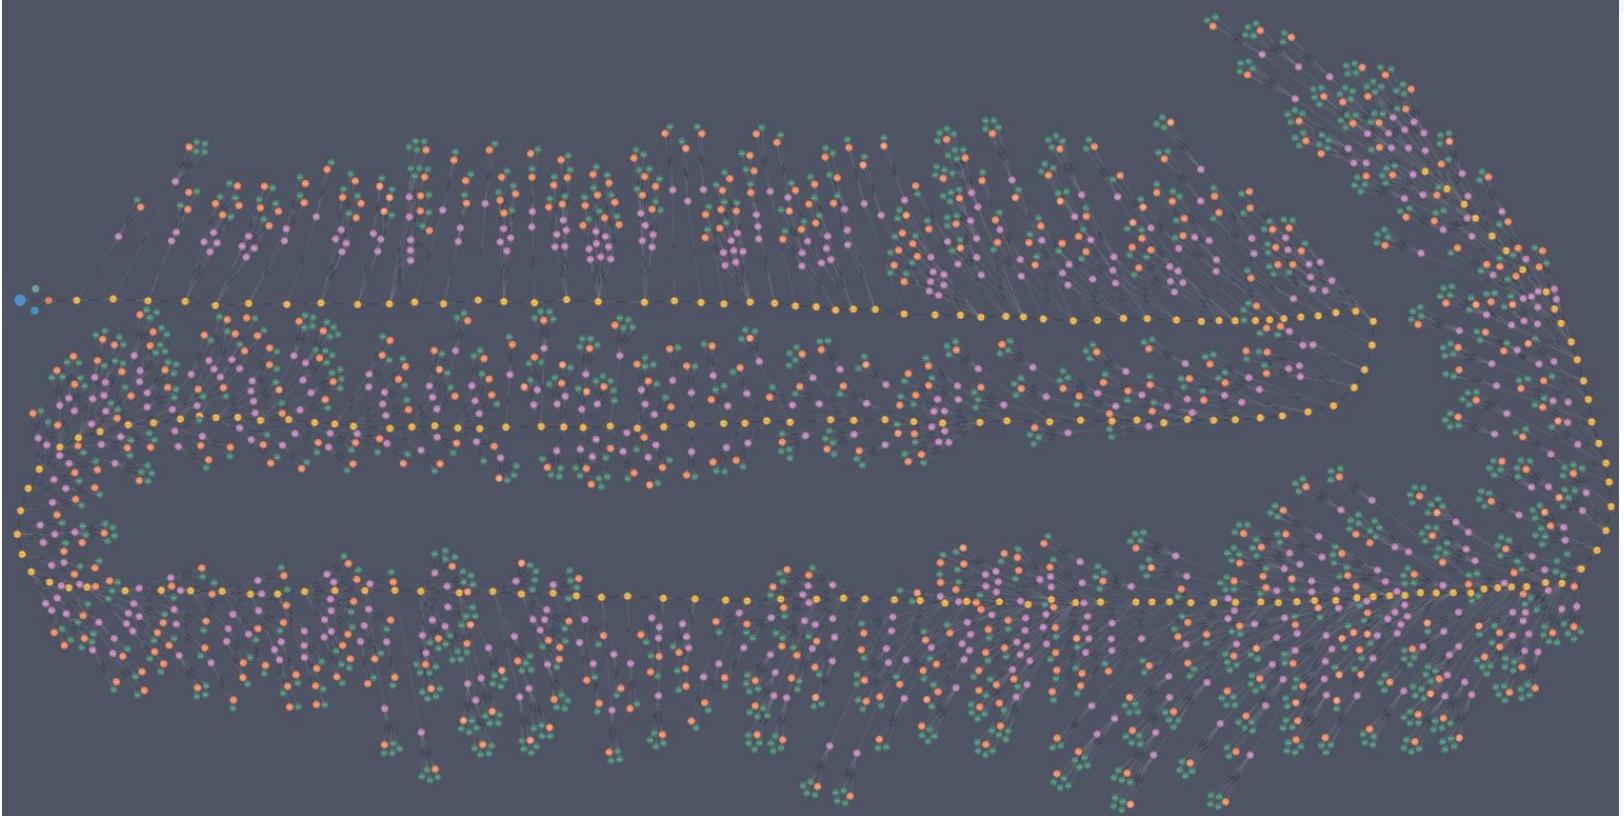
This is a Graph ...



This is a Graph, too ...

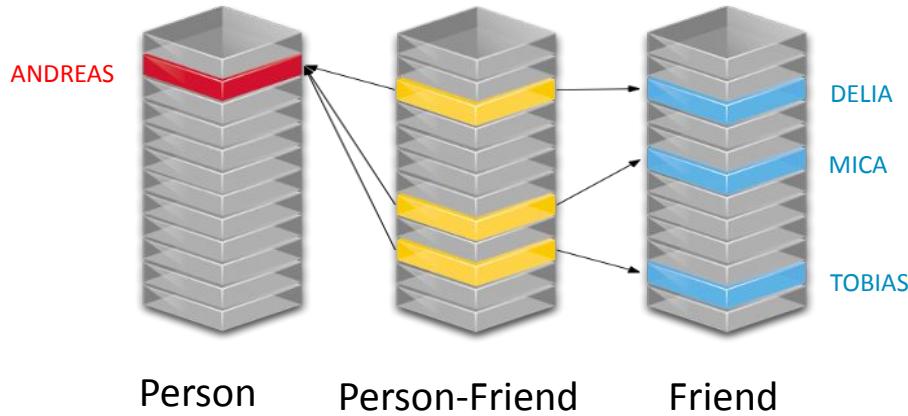


... and this is a Graph*! Guess, what graph could it be?

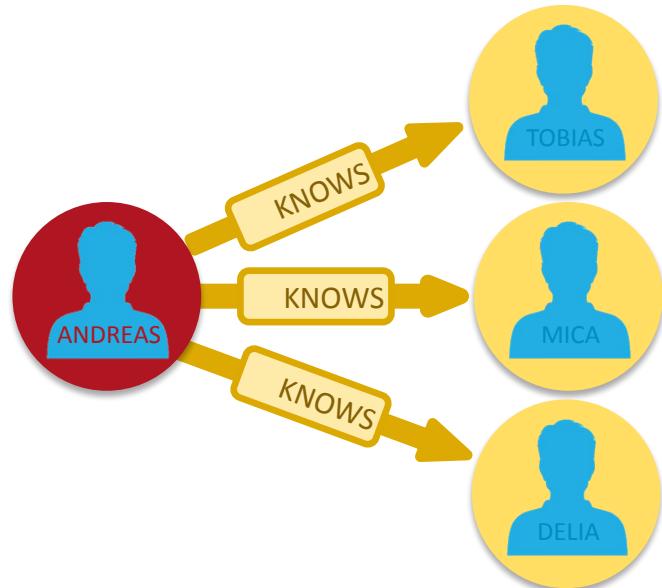


RDBMS vs Graph Model

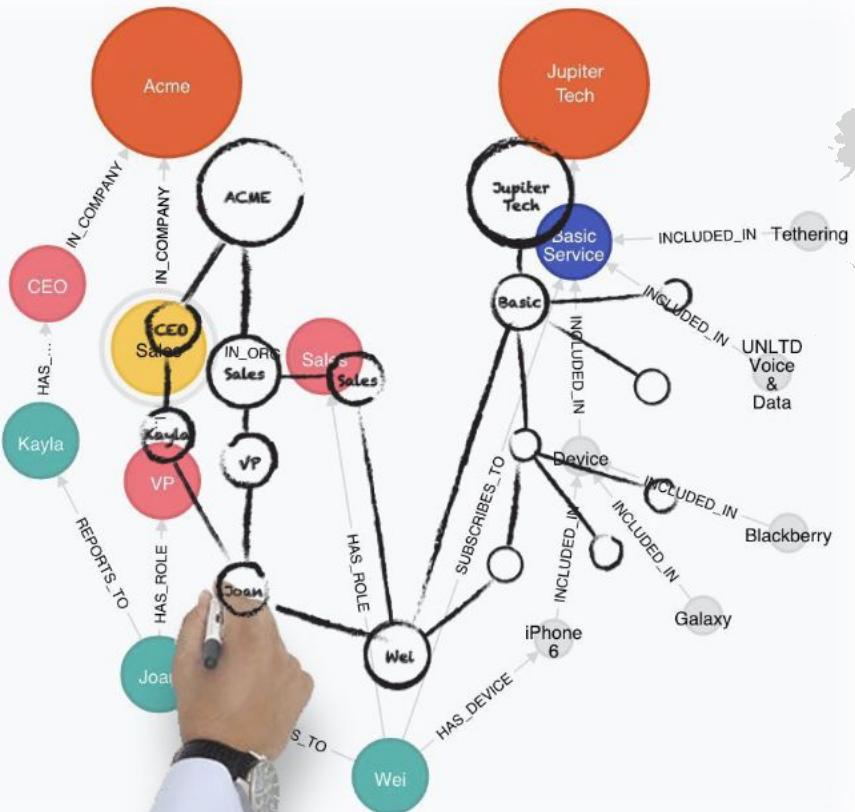
Relational Model



Graph Model



The Whiteboard Model Is the Physical Model



Labeled Property Graph Model

Nodes

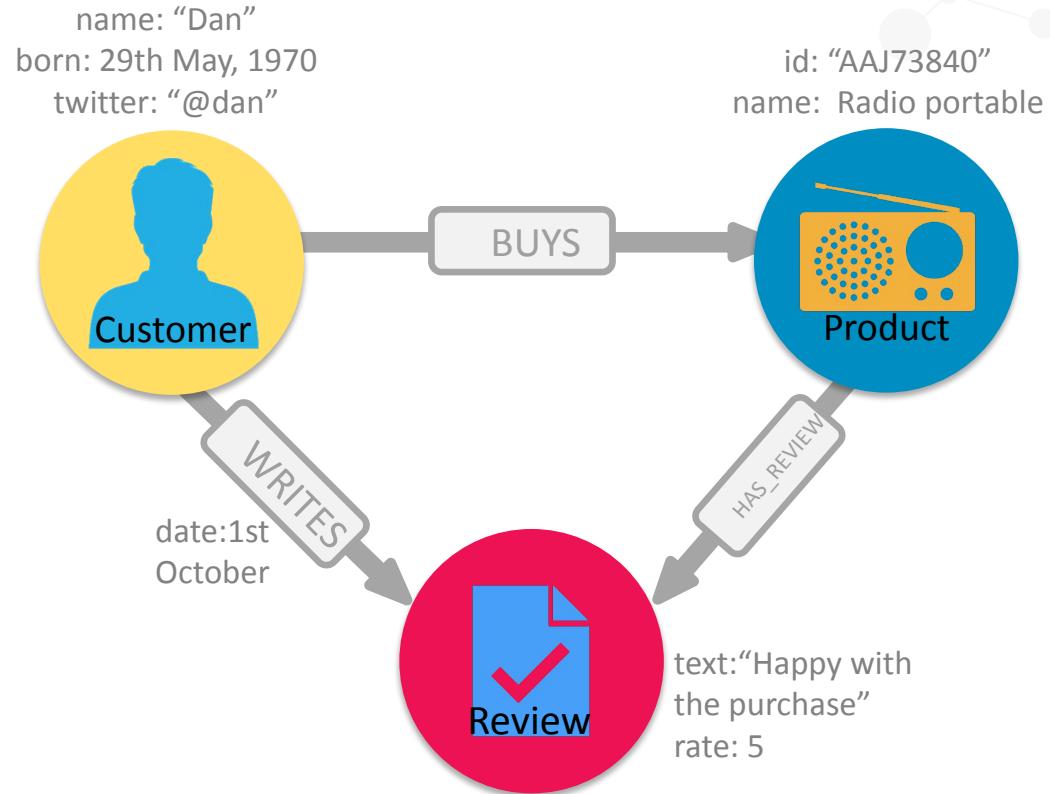
- Can have *Labels* to classify nodes
- Labels have *native indexes*

Relationships

- Relate nodes by type and direction

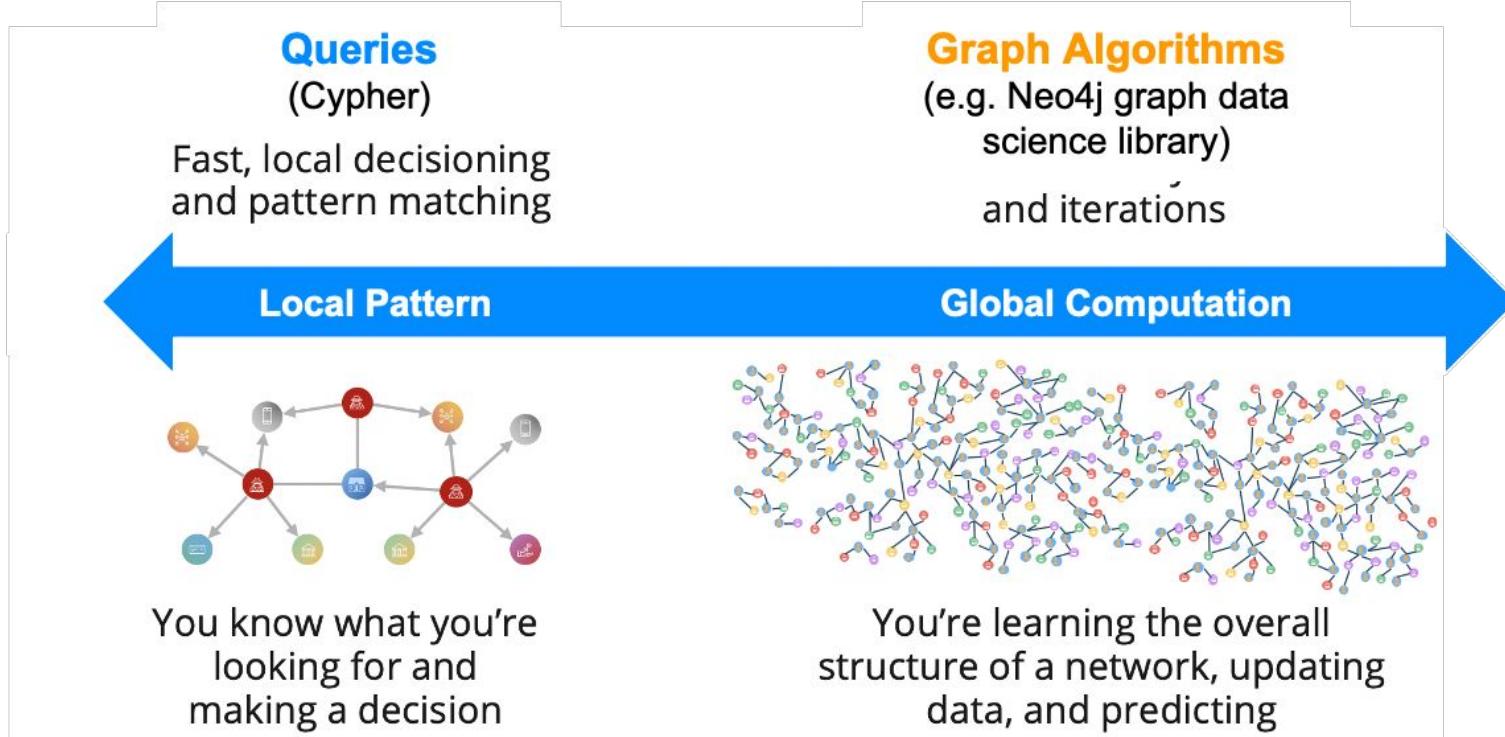
Properties

- Attributes of Nodes & Relationships
- Stored as Name/Value pairs
- Can have indexes and composite indexes



What is Graph Data Science

Data science when *relationships matter*

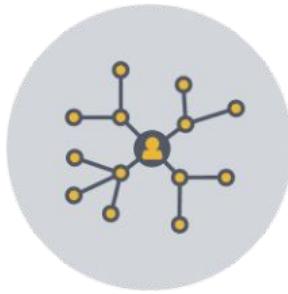


Graph Algorithms Verticals



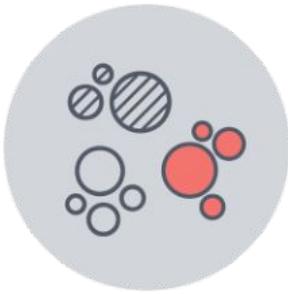
Pathfinding and Search

Finds optimal paths or evaluates route availability and quality.



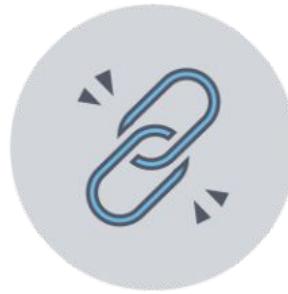
Centrality

Determines the importance of distinct nodes in the network.



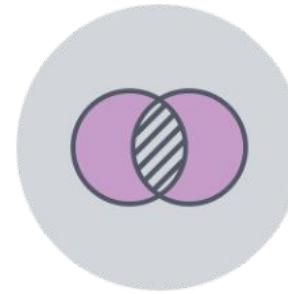
Community Detection

Detects group clustering or partition.



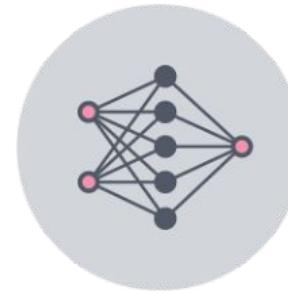
Heuristic Link Prediction

Estimates the likelihood of nodes forming a future relationship.



Similarity

Evaluates how alike nodes are by neighbors and relationships.



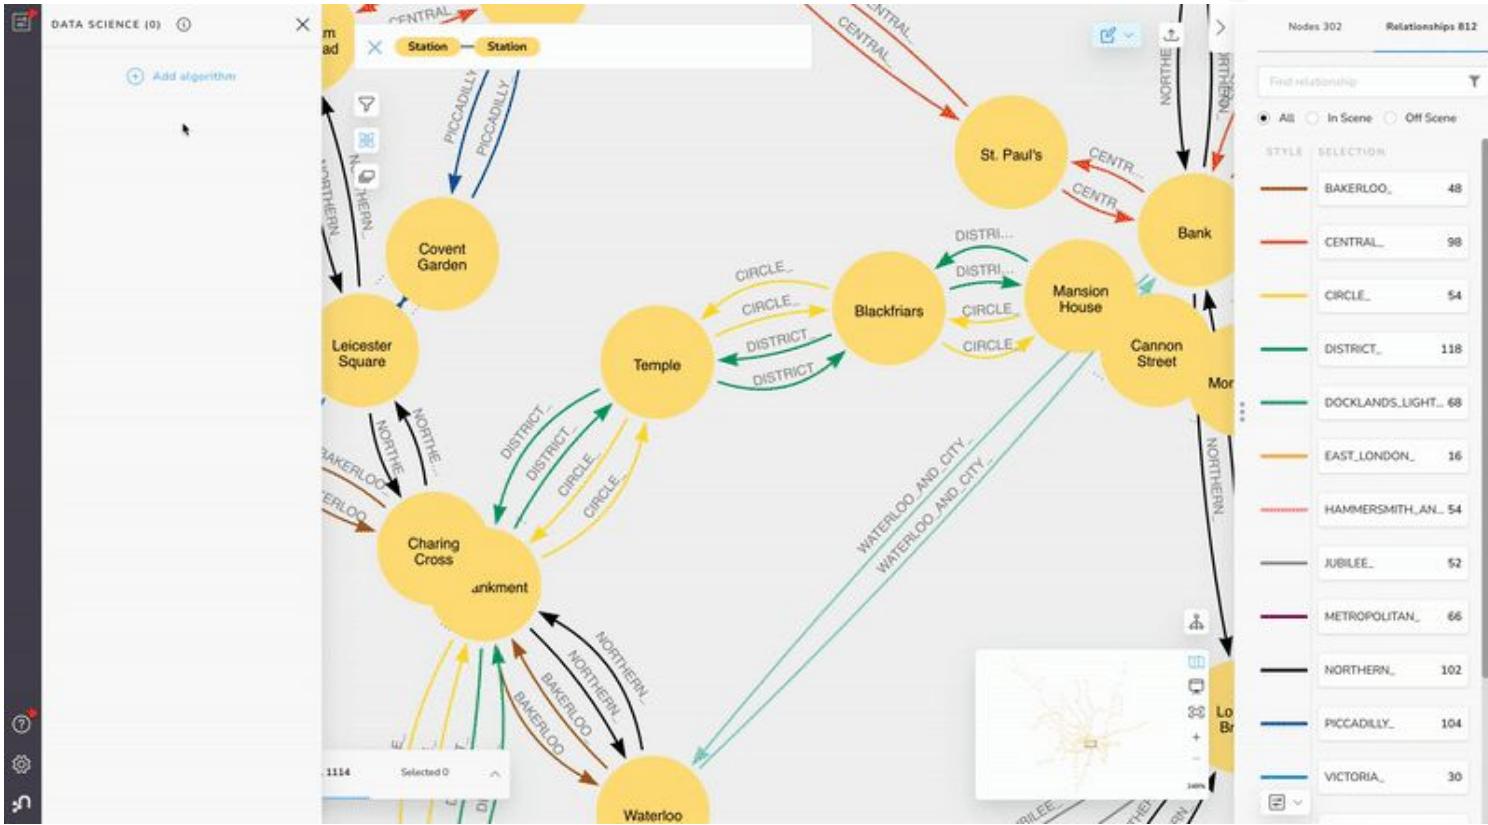
Embeddings

Learns graph topology to reduce dimensionality for ML

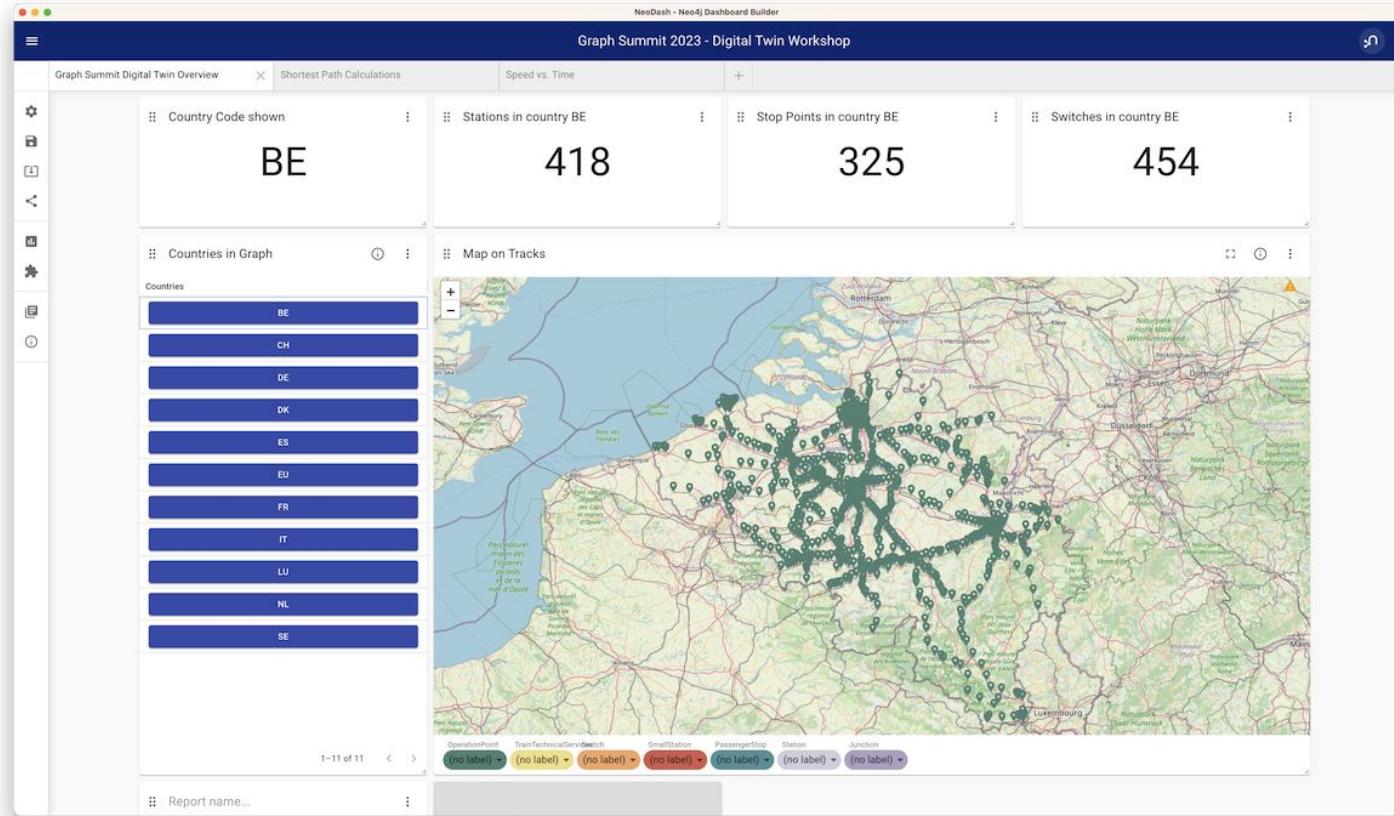


65+ graph algorithms in Neo4j

Neo4j Visualisation Platform - Neo4j Bloom



NeoDash - Dash Boarding with Graph Data



Use Case Explanation

Digital Twin - An Overview



What is a Digital Twin?



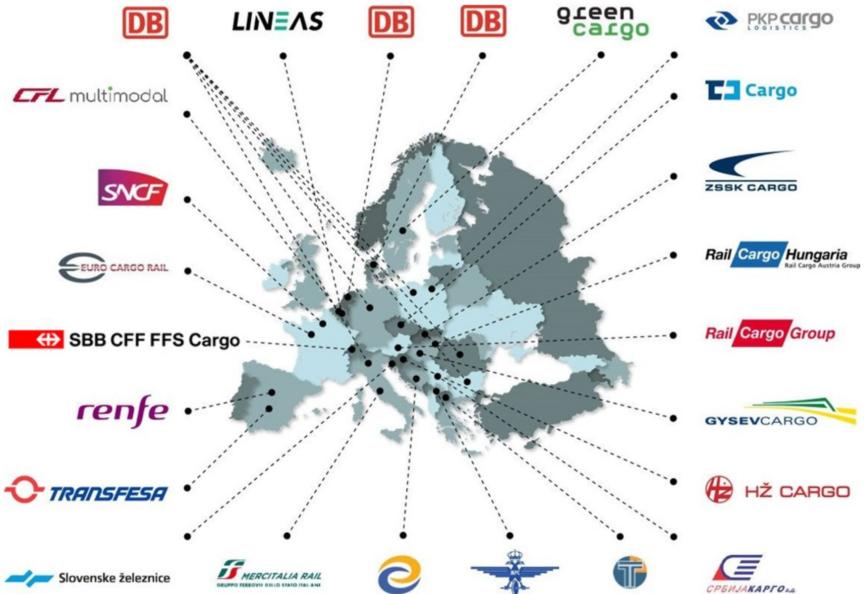
A **Digital Twin** is a digital representation of an intended or actual real-world physical product, system, or process (a *physical twin*) that serves as the effectively indistinguishable digital counterpart of it for practical purposes, such as **simulation, integration, testing, monitoring and maintenance.**



It has been done before

EU Railway Network - Track & Trace

- Challenge: **Legacy technology** could not track and analyze train journeys
- Solution: **Neo4j Knowledge Graph**
- Identify and avoid bottlenecks (DB)**



Why do we need a Digital Twin (some reasons)

Improved efficiency

It can optimize its operations and reduce costs by simulating different scenarios and making data-driven decisions.



Enhanced safety

It can identify potential hazards and test safety measures to improve safety for passengers and employees.



Predictive maintenance

It can monitor asset condition in real-time, predict maintenance needs, and increase asset lifespan.



Improved customer experience

A digital twin can simulate disruptions and help proactively address issues to enhance the customer experience and increase satisfaction.



POIs included



Unleash the power of the graph



Abstracting from this use case

Digital Twin and Networks are everywhere

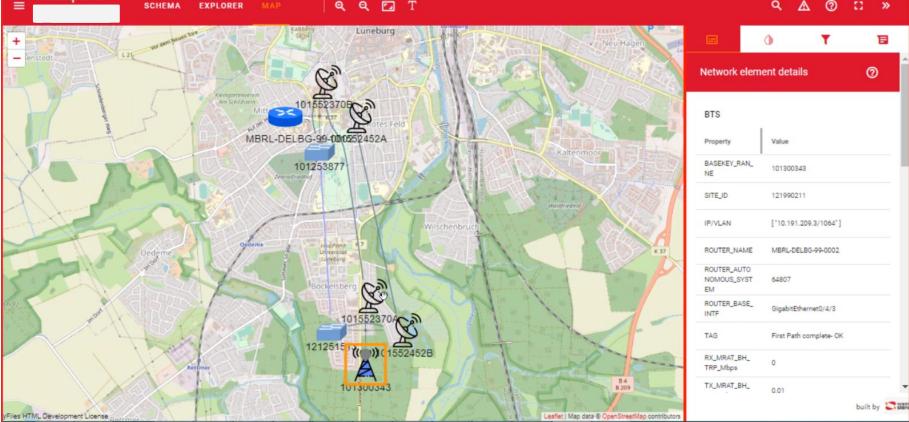
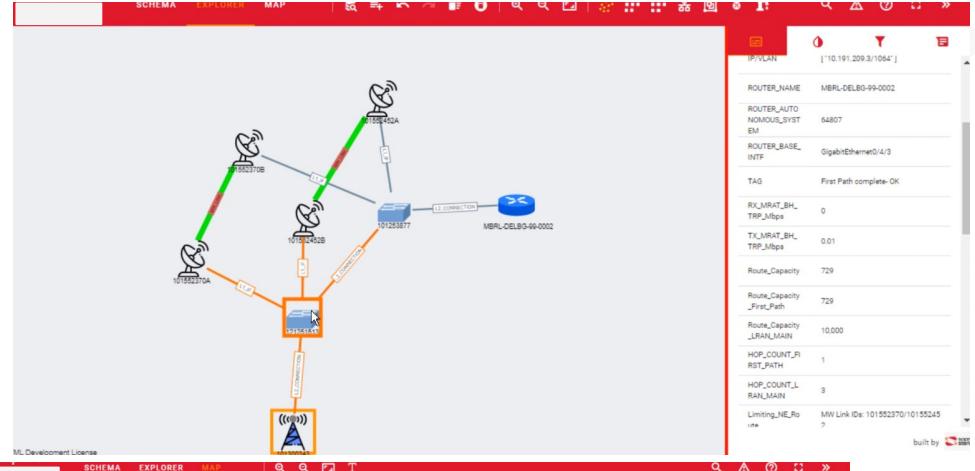


Many other use cases are networks, too!

- Supply Chain
- Mobile phone networks
- Power Grids / Gas Grids / Fibre Networks
- Social Networks
- Patient Journeys
- etc.

Mobile Network Operations - Large Mobile Provider

- Digital Twin of Mobile Network
- Planning and operation
 - Low signal strength of tower
 - Moving equipment around
- Simulation of possible failure scenarios
- Repair & maintenance support
 - Supply Chain Impact





“We wanted to create a solution that exploits artificial intelligence, integrates current and historical AIS positions as well as multiple data feeds and APIs. Such an effort would require a world-class infrastructure. That’s why we selected Neo4j.”

David Levy
Chief Marketing Officer
OrbitMI

Customer Case Study: Logistics and Supply Chain

Plan maritime routes based on distances, costs, and internal logic.

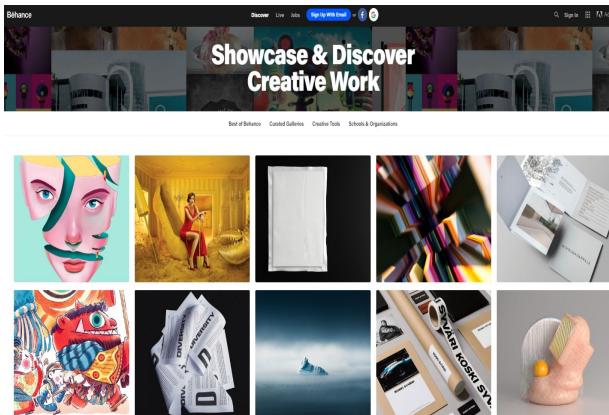
Results:

- **Subsecond** maritime routes planning
- Reduce global carbon emissions **60,000 tons**
- **12-16M ROI** for OrbitMI customers



Adobe Behance

Social Network of 10M Graphic Artists



Background

- Social network of 10M graphic artists
- Peer-to-peer evaluation of art and works-in-progress
- Job sourcing site for creatives
- Massive, millions of updates (reads & writes) to Activity Feed
- *150 Mongos to 48 Cassandras to 3 Neo4j's!*

Business Problem

- Artists subscribe, appreciate and curate “galleries” of works of their own and from other artists
- Activities Feed is how everyone receives updates
- 1st implementation was 150 MongoDB instances
- 2nd implementation shrunk to 48 Cassandras, but it was still too slow and required heavy IT overhead

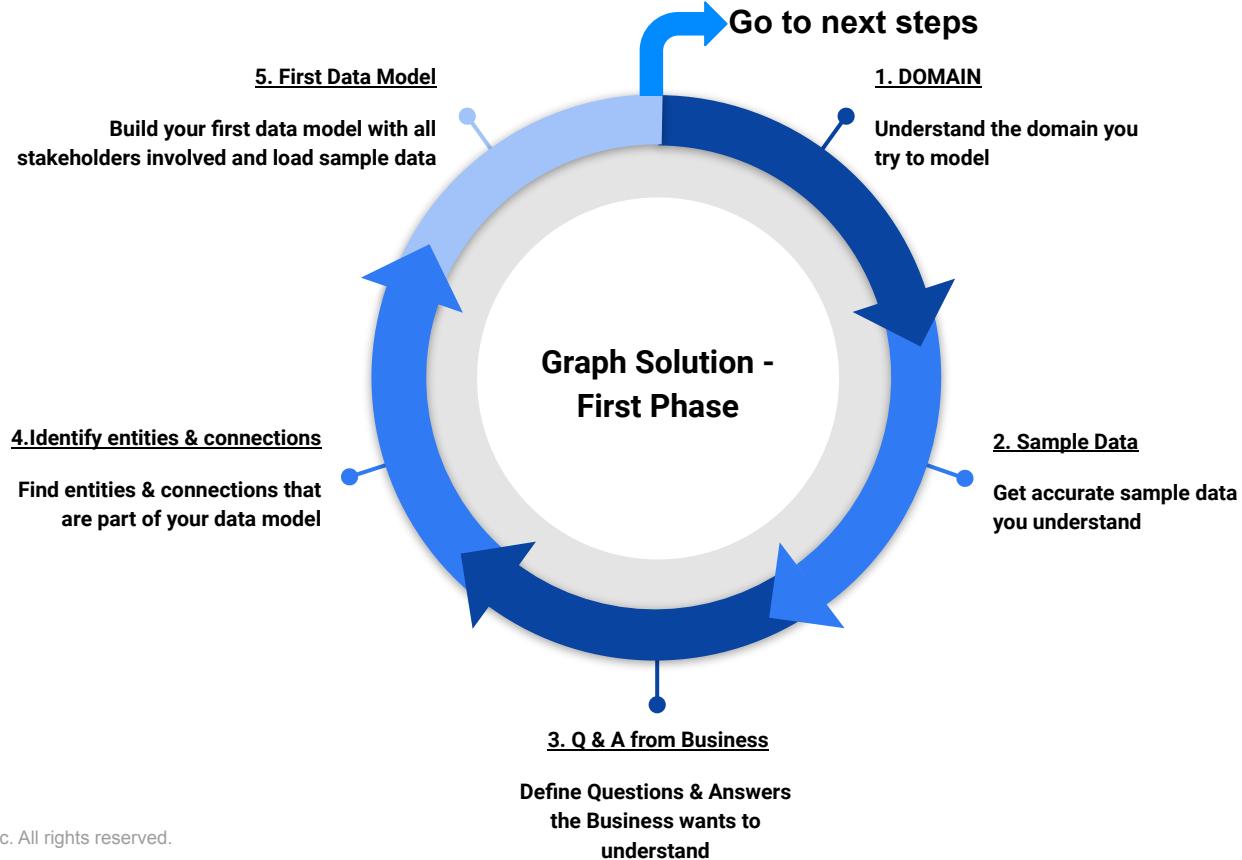
Solution and Benefits

- 3rd implementation shrunk to 3 Neo4j instances
- Saved over \$500k in annual AWS fees
- Reduced data footprint from 50TB to 40GB
- Significantly easier to introduce new features like, “New projects in your Network”

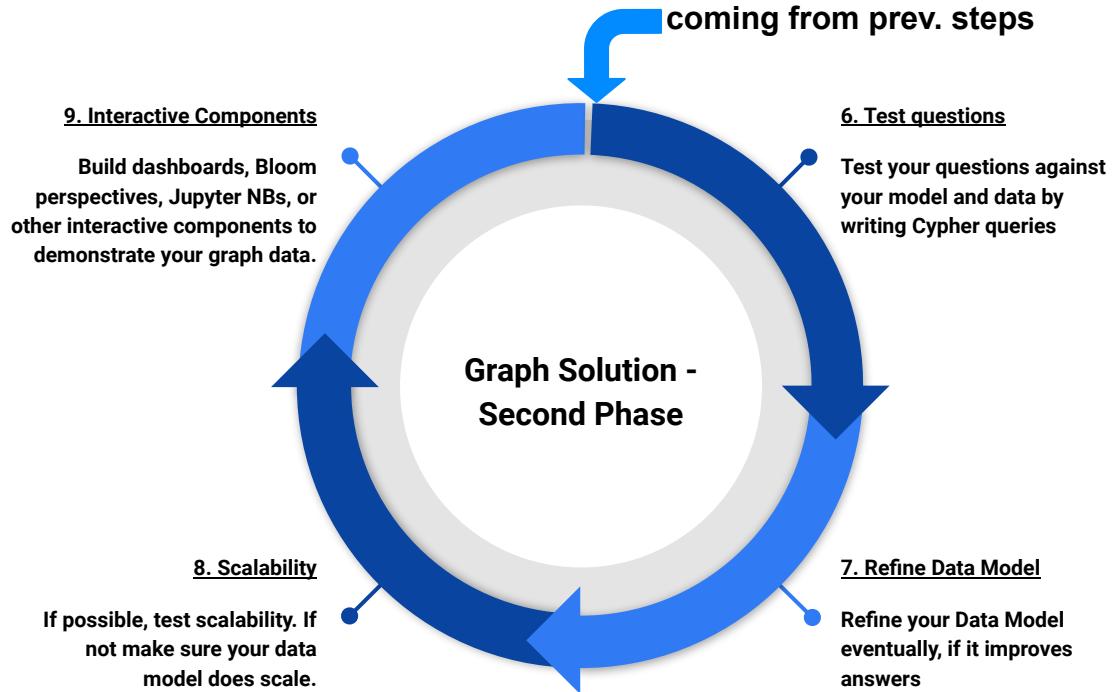
Building our solution



High Level Approach building a Graph Solution



High Level Approach building a Graph Solution



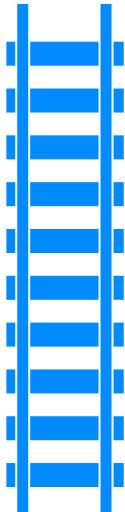
Stakeholders recommended to build a graph solution

<u>Domain Experts</u>	<u>Consultants / Developers</u>
	<ul style="list-style-type: none">• Add domain knowledge
	<ul style="list-style-type: none">• Translate questions into queries / scripts
	<ul style="list-style-type: none">• Provide questions they want to ask
	<ul style="list-style-type: none">• Build the graph
	<ul style="list-style-type: none">• Provide answers and rating for results to above questions
	<ul style="list-style-type: none">• Build and operate data loading (ETL process)
	<ul style="list-style-type: none">• Know what is missing today
	<ul style="list-style-type: none">• Build UIs, Dashboards, etc.
	<ul style="list-style-type: none">• Help to precise data model objects like labels, relationships, etc.
	<ul style="list-style-type: none">• Maintain / extend graph

The Sample Data Set



Sample Data: Railway Network Digital Twin dataset



Provided Data:

- Consists of a Railway Network + **Operation Point (OP)**
- **It includes attributes like:**
 - Sections,
 - Stations and other OPs,
 - Geo location information,
 - **Section length and Speed,**
 - other parameters
- Add on from us: **Point of Interest (POI)** along the network
- **Format:** CSV → Generated from XML
- Origin of Data: <https://data-interop.era.europa.eu/search>

What Does the Sample Data Looks Like?



- Operation Points Data*:
- Sections & Speed Data*:
- Point-of-Interest Data:
 - <https://github.com/neo4j-field/gsummit2023/tree/main/data>

Origin of data:

- <https://data-interop.era.europa.eu/search>

The data set is public available and can be used under the license and conditions mentioned by OpenDB. We have translated the data to make it better understandable to everyone outside Germany.

The Data Explained - In an Easy Way

“The data can be seen as a highway that has ramps throughout its entire path.

Each ramp can be an **Operation Point** (OP) like a Station, a Switch, etc.

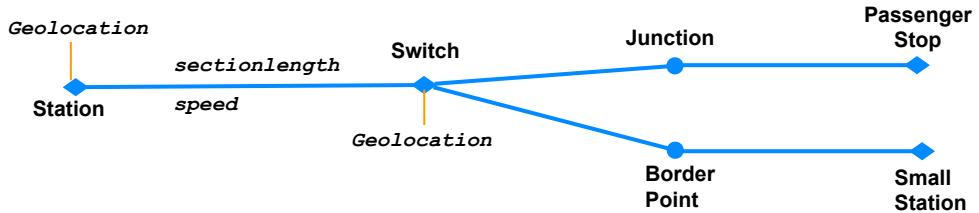
Sections have a length and a speed associated. **Sections have a start point and an endpoint and will be inter-connecting Operation Points.**

An Operation Point has a name and is referred to by a relationship with a country code.

Operation Points (OP) - Data Explanation (Nodes)

CSV Header titles:

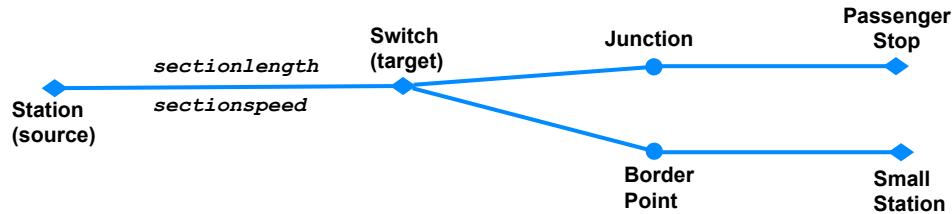
- **id**: the internal number of the OP
- **extralabel**: the kind of OP we deal with, e.g. Station, Junction, Switch, etc.
- **name**: the name of a OP
- **latitude**: of the OP
- **longitude**: of the OP



Section Connection Data Explanation (Relationships)

CSV Header titles:

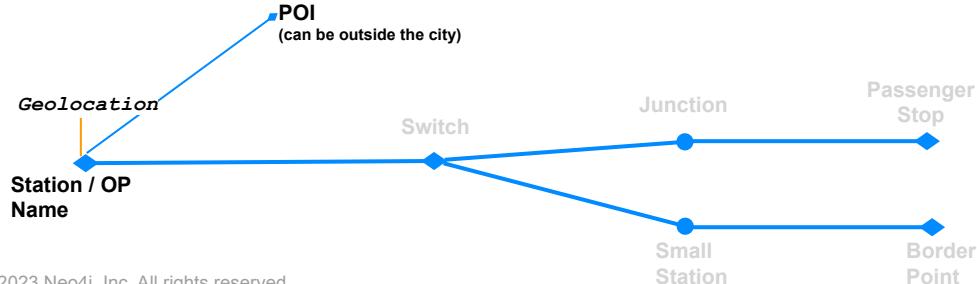
- **source**: start OP for this section
- **target**: end OP for this section
- **sectionlength**: the length in km of that section
- **speed**: max speed allowed on that section



Point of Interest (POI) Data Explanation

CSV Header titles:

- **CITY:** City the POI is in or close by
- **POI_DESCRIPTION:** A short description of the POI
- **LINK_FOTO:** is a foto of the POI
- **LINK_WEBSITE:** is the link of the website corresponding to POI
- **LAT:** Latitude of the POI
- **LONG:** Longitude of the POI
- **SECRET:** True if not well known, False if well known (e.g. Berlin)



Data Modeling - Writing down Questions & Answers



Data Modeling - The Questions are the Input

1. How long is the **journey** from **station X** to station of **station Y**?
2. Do I have alternative routes if a **station** in my journey is down?
3. What is the *shortest distance/path*, between **stations X** and **Y**?
4. Can I **geolocate** stations in my railway network?
5. What are the stations expecting the most traffic?
6. What **POIs has** a station (city) along my route?

Note: Aggregations are also possible, but we focus on the “graphy” queries here!

Expected Answers



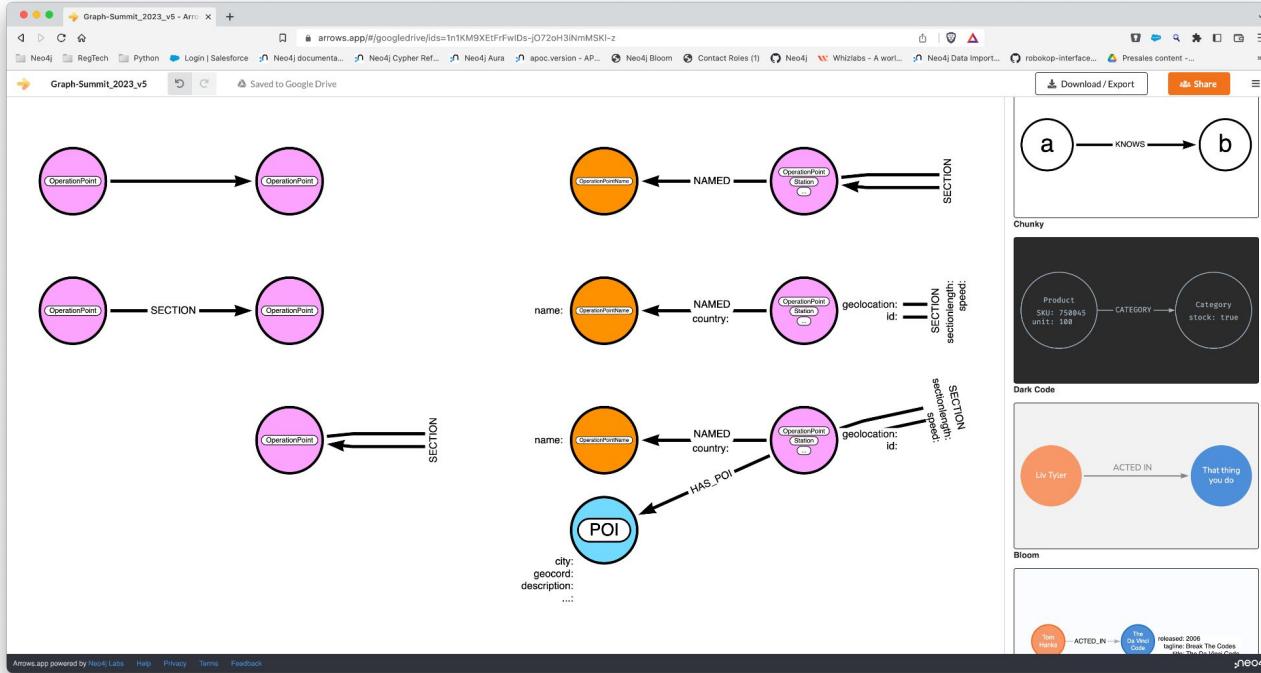
1. Your journey is 30 stations and 239km long
2. Your alternative route is 36 stations and 271km long
3. The shortest route between X and Y is XXX km
4. The geo location of your station is Lat / Long
5. The station with the most traffic is ... ?
6. The following POIs are along your route ...

Data Modeling - Building the Model



Tools for Graph Data Modeling

Arrows Tool → <https://arrows.app>



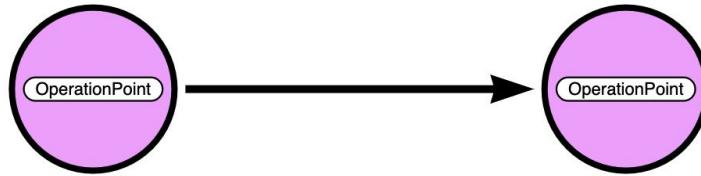
Data Modeling Using the Questions as Input

Remember those questions?

1. How long is the **journey from station X to** station of **station Y**?
2. Do I have alternative routes if a **station** in my journey is down?
3. What is the **shortest distance/path, between stations** X and Y?
4. Can I **geolocate** stations in my rail **railway** network?
5. What are the stations expecting the most traffic?
6. What **POIs has** a station (city) along my route?

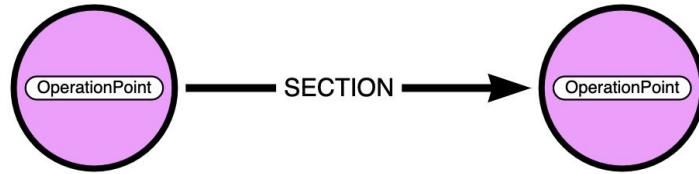
Building a First Data Model 1/6

- Nouns in your questions are the nodes, verbs are the relationships
- Nodes have a **Label**, Relationships MUST have a **Type**
- Properties should help to uniquely identify Nodes and/or answer questions



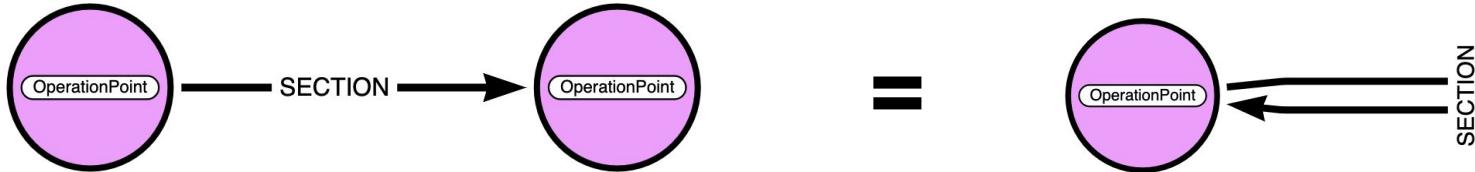
Building a First Data Model 2/6

- Nouns in your questions are the nodes, verbs are the relationships
- Nodes have a **Label**, Relationships MUST have a **Type**
- Properties should help to uniquely identify Nodes and/or answer questions



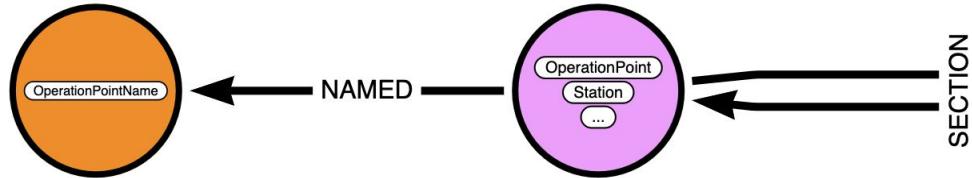
Building a First Data Model 3/6

- Nouns in your questions are the nodes, verbs are the relationships
- Nodes have a **Label**, Relationships MUST have a **Type**
- Properties should help to uniquely identify Nodes and/or answer questions



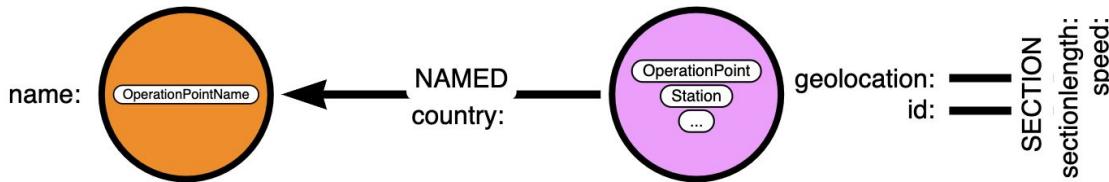
Building a First Data Model 4/6

- Nouns in your questions are the nodes, verbs are the relationships
- Nodes have a **Label**, Relationships MUST have a **Type**
- Properties should help to uniquely identify Nodes and/or answer questions



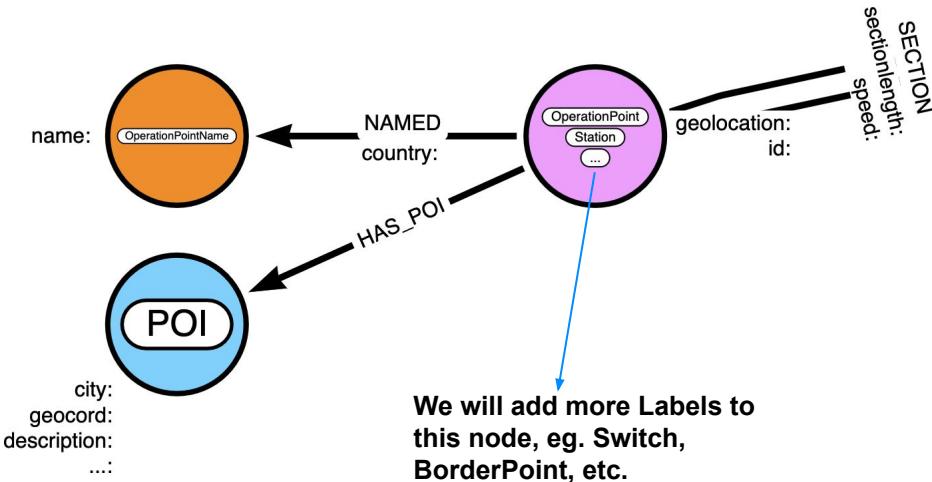
Building a First Data Model 5/6

- Nouns in your questions are the nodes, verbs are the relationships
- Nodes have a **Label**, Relationships MUST have a **Type**
- Properties should help to uniquely identify Nodes and/or answer questions



Building a First Data Model 6/6

- Nouns in your questions are the nodes, verbs are the relationships
- Nodes have a **Label**, Relationships MUST have a **Type**
- Properties should help to uniquely identify Nodes and/or answer questions

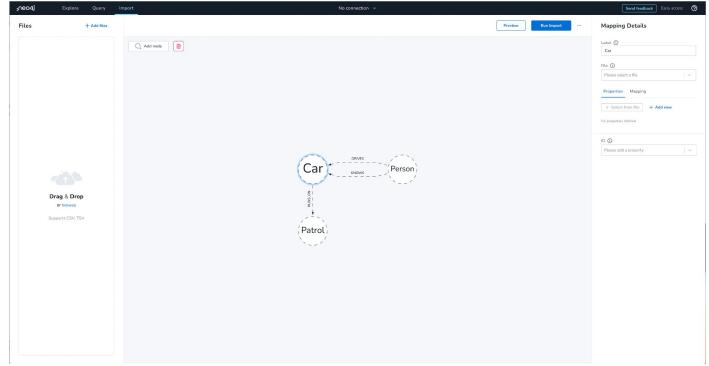


Data Loading



Data Loading - How Can I Do That?

- Cypher CSV Load
- APOC Library
- neo4j admin import command
- ETL Tool
- Python API
- Neo4j Workspace*
- 3rd Party Tools (Apache Airflow or HOP, Cloud provider tools, etc.)



[Neo4j Workspace*](#)

More about [Data loading, etc.](#) on the [Neo4j Website](#)

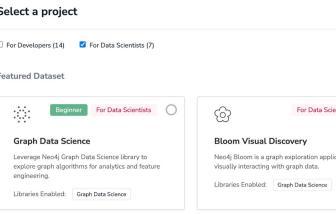
Workshop Time

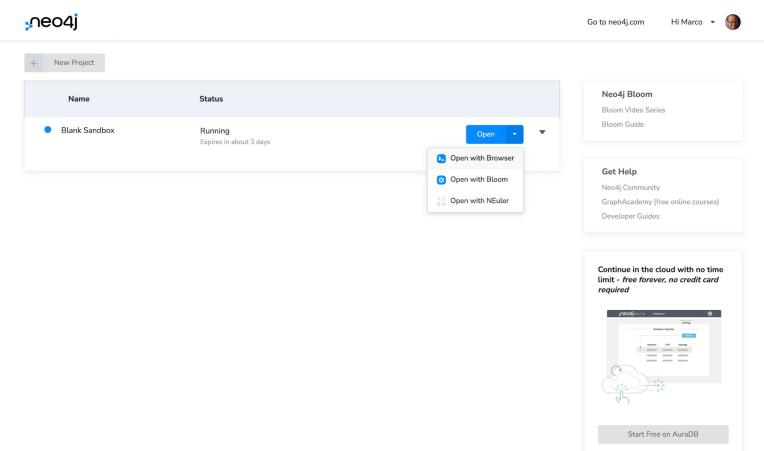
Let's build the solution together ...



Lets do it together ...

1. Open the Github page at: <https://github.com/neo4j-field/gsummit2023>
2. Open <https://sandbox.neo4j.com/> and create a Neo4j Sandbox (Google login required!)

a) 

b) 

3. Open Neo4j Browser and login with your credentials
4. Later on open NeoDash at: <http://neodash.graphapp.io/>

Improving the Solution further

To add even more value to your business ...



How could this “Mini Digital Twin” be enhanced?

Extending the Graph is easy, e.g. with further section information:

- For example with possible **pace on section**, **section quality metrics**, **mobile network coverage**, etc.
- Sensor information about switches to quickly find problems and re-route traffic
- The more related information is added, the more use cases could benefit from the digital twin

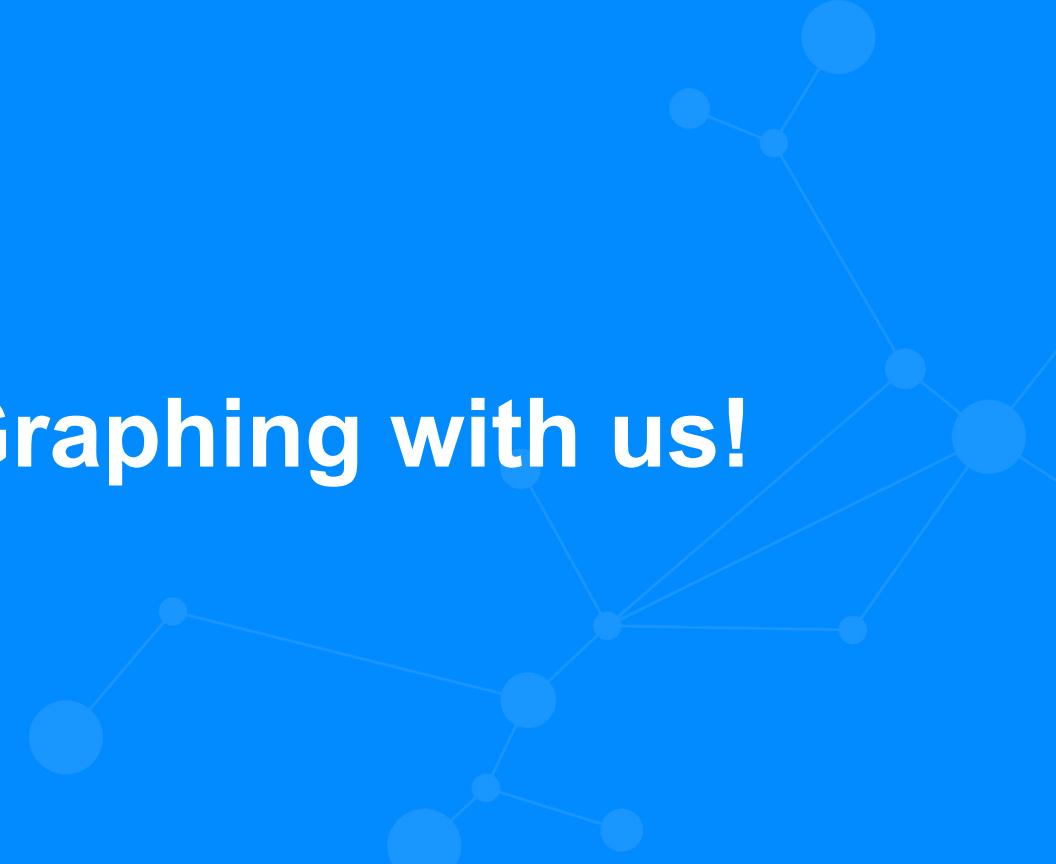
Going even further:

- Additional information about **trains running on the sections** enable additional benefit:
 - Section and schedule optimization
 - Simulation of new sections (to keep up with growing traffic)

Q & A



Thanks for Graphing with us!



Contact us:

Marco.Deluca@neo4j.com

Luis.Salvador@neo4j.com

Gal.Bello@neo4j.com

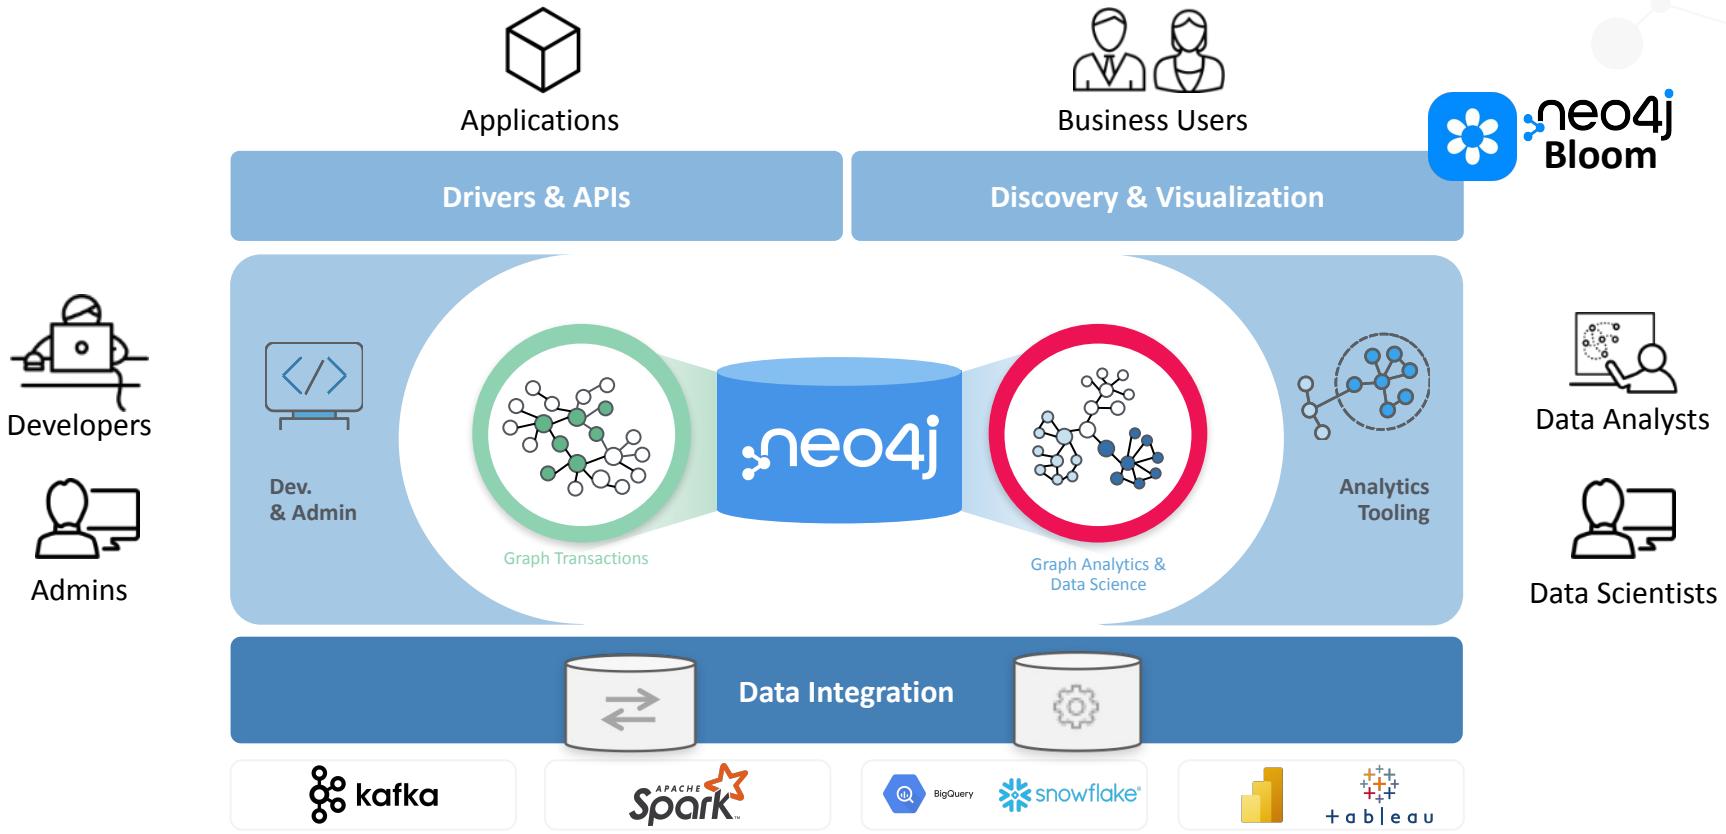


Contact us at
sales@neo4j.com

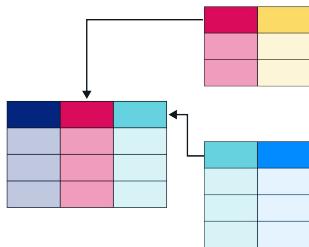
Additional Information



Neo4j Graph Data Platform Overview



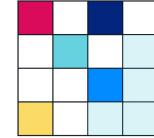
Use context to unlock insights



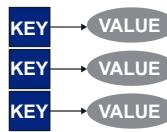
DOCUMENT

```
{
  "book_id": 1,
  "title": "Graph Databases",
  "release_date": "2013-06-01T1",
  "authors": [
    {
      "name": "Emil Eifrem"
    },
    {
      "name": "Jim Webber"
    },
    {
      "name": "Alan Hay"
    }
  ]
}
```

COLUMN



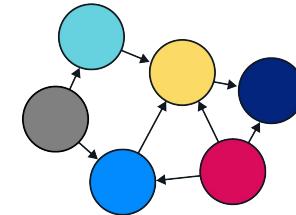
KEY-VALUE



Relational Databases

Don't handle relationships well

Practically limited to three degrees of separation



Other NoSQL Databases

Don't handle relationships at all

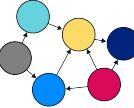
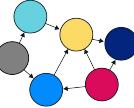
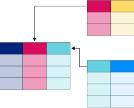
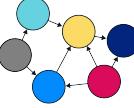
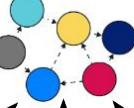
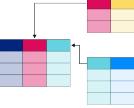
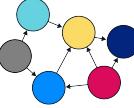
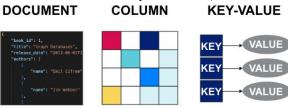
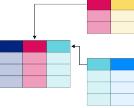
No language or structures for handling relationships: joins done through the application

Graph Databases

Natively store & query relationships

Queries run 1000x faster at scale, up to 1000+ hops

How Neo4j differentiates from other databases

	neo4j Native Graph DB	Non-Native Graph DB	RDBMS
Visualization			
Queries	Cypher (graphs)-[are]->(everywhere)	openCypher, Gremlin, Proprietary	SQL
Processing			
Storage		<p>DOCUMENT COLUMN KEY-VALUE</p> 	

Cypher

```

MATCH (u:Customer
{customer_id:'customer-one'})-[:BOUGHT]->(p:Product)<-
[:BOUGHT]-(peer:Customer)-[:BOUGHT]->(reco:Product)
WHERE not (u)-[:BOUGHT]->(reco)
RETURN reco as Recommendation, count(*) as Frequency
ORDER BY Frequency DESC LIMIT 5;
  
```

SQL

```

SELECT product.product_name AS Recommendation,
count(1) AS Frequency
FROM product, customer_product_mapping, (SELECT
cpm3.product_id, cpm3.customer_id
FROM Customer_product_mapping cpm,
Customer_product_mapping cpm2,
Customer_product_mapping cpm3
WHERE cpm.customer_id = 'customer-one'
AND cpm.product_id = cpm2.product_id
AND cpm2.customer_id != 'customer-one'
AND cpm3.customer_id = cpm2.customer_id
AND cpm3.product_id NOT IN (SELECT DISTINCT
product_id
FROM Customer_product_mapping cpm
WHERE cpm.customer_id = 'customer-one')
) recommended_products
WHERE customer_product_mapping.product_id =
product.product_id
AND customer_product_mapping.product_id IN
recommended_products.product_id
AND customer_product_mapping.customer_id =
recommended_products.customer_id
GROUP BY product.product_name
ORDER BY Frequency DESC
  
```

How to deploy Neo4j - 3 deployment models available

DB-as-a-Ser vice



Fully-managed SaaS
Consumption-based pricing

Cloud-native
Self-service deployment
No access to underlying
infrastructure and systems

Cloud Managed Services



White-glove managed service
by Neo4j experts

Fully customizable deployment
model and service levels
Operate In own data centers
or Virtual Private Cloud

Self-hosted



For private, hybrid or
lift-and-shift cloud
Bring-your-own-license

Full control of your environment
Run in any cloud, in your account